

# Software Architecture and Design Specification (SAD)

---

Project: API Rate Limiter

Team mem: Ankita muni PES1UG23CS081

Version: 1.0

Arya S PES1UG23CS107

Authors: QuadCore

Manasa S PES1UG24CS811

Date: 18-09-2025

Likitha H PES1UG24CS810

Status: Draft

## 1. Introduction

### 1.1 Purpose

This document specifies the architecture and design of the API Rate Limiter system. The purpose is to ensure backend services are protected from abuse, denial-of-service attacks, and excessive API consumption, while enforcing fairness among clients.

### 1.2 Scope

Covers API request tracking, limit enforcement, error handling, logging, monitoring, and configuration. Excludes backend API business logic.

### 1.3 Audience

DevOps Engineers, Frontend and Backend Developers, API Consumers, Security Analysts, QA/ Test Engineers

### 1.4 Definitions

API: Application Programming Interface

Rate Limiting – Restricting API calls per time unit

Token Bucket/Leaky Bucket – Algorithms for rate control

HTTP 429 – Standard error code for “Too Many Requests”

Retry-After – Header to indicate cooldown time

Redis – In-memory store for counters

## 2. Document Overview

### 2.1 How to use this document

This SAD provides a complete architectural view of the API Rate Limiter, including goals, component diagrams, sequence flows, API design, error handling, and security architecture.

## 2.2 Related Documents

- SRS v1.0 – Software Requirements Specification
- STP v1.0 – Software Test Plan
- RTM – Requirements Traceability Matrix

## 3. Architecture

### 3.1 Goals & Constraints

Goals:

- Enforce configurable rate limits
- Ensure <1 ms latency overhead per request (p99)
- Achieve 99.99% uptime
- Provide fair usage and prevent abuse


Constraints:

- Stateless servers with distributed counters
- Must scale horizontally
- must integrate with REST APIs
- handle 1M+ requests/day

### 3.2 Stakeholders & Concerns

API Clients, Admins/DevOps, Security Teams, Developers

### 3.3 Component (UML) Diagram



### 3.4 Component Descriptions

- Request Interceptor: Sits in front of APIs, checks request counts.
- Policy Manager: Implements algorithms (fixed, sliding, token bucket).
- Counter Store: Distributed backend (Redis) for consistent state.
- Admin API: REST endpoints for configuring limits and viewing stats.
- Logger/Monitor: Sends structured logs to monitoring tools.

### **3.5 Chosen Architecture Pattern and Rationale**

A Layered + Middleware architecture is chosen. This ensures clear separation: interception, policy enforcement, storage, and monitoring. Microservices style is avoided to keep latency minimal.

### **3.6 Technology Stack & Data Stores**

Languages/Frameworks: Node.js / Java Spring Boot / Python Flask

Data Store: Redis or Memcached for counters

Protocols: HTTP/HTTPS (REST)

Monitoring: Prometheus

### **3.7 Risks & Mitigations**

- Redis downtime → fallback to local memory + auto-retry
- High load spikes → burst handling + horizontal scaling
- Config tampering → secured RBAC for Admin API
- Log overload → structured logging with sampling

### **3.8 Traceability to Requirements**

RL-F-001 (Enforce rate limit) → Request Interceptor + Counter Store

RL-F-004 (Admin API) → Admin Module

RL-NF-005 (Latency ≤1ms) → System-wide optimization

### **3.9 Security Architecture (STRIDE)**

Spoofing → Authenticated clients, TLS

Tampering → Signed configs, role-based access

Repudiation → Audit logs with timestamps

Info Disclosure → No sensitive logs, TLS enforced

DoS → Rate-limiting policies, burst handling


Elevation of Privilege → RBAC for admin endpoints

## **4. Design**

### **4.1 Design Overview**

The system intercepts requests, checks counters in Redis, applies policies, and either forwards the request or blocks with HTTP 429.

## 4.2 UML Sequence Diagrams



## 4.3 API Design

Endpoint: /check-limit

Method: GET

Response: {status: allowed/blocked, remaining: int, reset\_time: timestamp}

Errors: 429 Too Many Requests

Endpoint: /admin/config

Method: POST

Request: {endpoint: '/api/orders', limit: 100, window: '1m'}

Response: {status: success}

Errors: 401 Unauthorized, 403 Forbidden

## 4.4 Error Handling, Logging & Monitoring

- Standardized JSON error responses.
- No sensitive info in logs.
- Logs include: client ID, timestamp, endpoint, status.
- Metrics: % requests blocked, average latency, top clients by usage.

## 4.5 UX Design

Client-facing: Clear error responses with retry-after

Admin-facing: REST APIs + optional dashboard for live statistics

## 4.6 Open Issues & Next Steps

Future support for adaptive rate limiting using ML models

Distributed policies across multi-cloud environments

Integration with API gateways (Kong, Nginx, Envoy)

## 5. Appendices

### 5.1 Glossary

API, Rate Limiting, HTTP 429, Token Bucket, Redis, RBAC

### 5.2 References

IEEE 42010

OWASP API Security Top 10

NIST SP 800-160

### 5.3 Tools

PlantUML, draw.io (for diagrams)

Swagger/OpenAPI (for API contracts)